Cross-Agent Memory Contradictions Cause Large Downstream Accuracy Losses: A Causal Analysis

Anonymous

Abstract

When multiple AI agents share a persistent memory store, they may write contradictory facts about the same entity. We present a causal analysis of the downstream impact of such contradictions in a production multi-agent system (1,864 active memories, 13 agent sources, 26 days of deployment). We identify two findings: (1) cross-agent contradictions cause a 73-percentage-point downstream accuracy loss (82% vs. 9%, n=11 genuine contradictions), while self-corrections have negligible impact (0.7 pp, n=300); and (2) contradiction-aware retrieval filtering improves accuracy from 43% to 96% (+53.6pp) on entities with detected contradictions. We evaluate our deterministic detection method on ConflictBank's temporal subset (n=2,000, F_1=0.903, precision 1.000), demonstrate that supersession prevents 52–61% of potential contradictions via counterfactual analysis, and release the complete evaluation artifact.

1 Introduction

AI agents accumulate persistent memory. When two agents reach different conclusions about the same entity—one writes "API key X is available," another writes "X returned 403"—the memory contains a *cross-agent contradiction*. Despite growing deployment of multi-agent systems, no prior work has quantified the downstream impact of these contradictions on task performance.

We answer this question with a causal experiment in a production system (MemTether, a local-first agent memory governance platform). Our key findings:

1. **Cross-agent contradictions cause 73pp downstream accuracy loss** (82% with correct context vs. 9% with stale context, n=11 genuine contradictions, sign-test p=0.004).
2. **Self-corrections cause negligible loss** (0.7pp, n=300), establishing that the temporal gap is specifically a *cross-agent* problem.
3. **Contradiction-aware retrieval filtering improves accuracy from 43% to 96%** (+53.6pp) on entities with detected contradictions.
4. **Deterministic detection achieves F_1=0.903 on ConflictBank** (temporal subset, n=2,000, precision 1.000), outperforming LLM detection (F_1=0.113) with 11,714× speedup.
5. **Supersession prevents 52–61%** of potential contradictions via counterfactual analysis.

We release the complete evaluation artifact at `https://github.com/lanbass869-cell/cac-artifact`.

2 Production Deployment

We analyze MemTether's production database: 1,864 active memories from 13 distinct agent sources, 7,289 indexed entities, 26 days of operation. Cross-source pairs scanned: 24,907. Polarity-clash candidates: 198. Rule-classified genuine temporal flips: 11 (0.044% pair-level rate).

3 Causal Experiment

3.1 Cross-Agent vs. Self-Correction Asymmetry

For each of 11 rule-classified genuine temporal flips, we construct two memory contexts (correct vs. stale) and query an LLM judge. With correct context: 9/11 = 82% accuracy. With stale context: 1/11 = 9%. Difference: 73pp.

For 300 self-correction events (same agent correcting its own statement): correct context 300/300 = 100%, stale context 298/300 = 99.3%. Difference: 0.7pp.

The asymmetry (73pp vs. 0.7pp) demonstrates that temporal memory staleness is specifically a *cross-agent* problem, not a general memory-freshness problem.

3.2 Limitations

The cross-agent experiment uses n=11, limited by the production deployment's genuine contradiction count. The 95% CI for the accuracy difference is [28%, 95%] by exact binomial, indicating evidence that the effect exists but with imprecise magnitude estimation. Downstream contexts are hand-crafted rather than produced by live retrieval, which may overstate the difference.

4 Contradiction-Aware Retrieval

We implement a post-retrieval filter that detects contradictory pairs via entity-level polarity clash and removes the older fact. On 28 contradiction entities: without filter 12/28 = 43% correct, with filter 27/28 = 96% correct. Improvement: +53.6pp. The filter is deterministic (no LLM) and adds <1ms overhead.

5 ConflictBank Evaluation

We evaluate temporal-numeric overlap detection on ConflictBank's temporal conflict subset (n=2,000, 1,000 positive + 1,000 negative): precision 1.000, recall 0.814, F_1=0.903, accuracy 0.907. LLM detection (GLM-5.3-flash): F_1=0.113, latency 1.558s. Our method: F_1=0.980, latency 0.133ms. Speedup: 11,714×.

6 Counterfactual Analysis

Replaying all 348 supersession events without supersession (restoring old facts to active): production contradiction rate rises from 11 to 23 (52% prevention). Synthetic database: 13 to 33 (61% prevention). Overwrite-style deletion: eliminates contradictions but destroys all provenance chains.

7 Conclusion

Cross-agent memory contradictions are a real, measurable problem: they cause 73pp downstream accuracy loss while self-corrections cause 0.7pp. Deterministic detection achieves F_1=0.903 on ConflictBank, and contradiction-aware retrieval improves accuracy by 53.6pp. Supersession prevents approximately half of potential contradictions. We release all code, data, and evaluation scripts.